Contents

Contributors:

- Daniel Fitz (Sanchez)

0.1 Assessment

- Final Exam (50%)
    - All goals tested except for developing applications
    - **Closed book exam**
- Two assignments
    - Individual programming assignment (25%)
        Implementation of a context-aware distributed application using RMI and "publish/subscribe" (event based architecture)
    - Group research/written assignment (25%)
        Group assignment on functionality and design issues of various distributed systems (e.g. grid computing, cloud computing, pervasive computing)

Chapter 1

Lecture Notes

1.1 Introduction to Distributed Systems

1.1.1 Definitions of Distributed Systems

> A collection of independent computers that appear to its users as a single coherent system (Andrew Tannenbaum)

> A system where I can't get my work done because a computer has failed that I've never even heard of (Leslie Lamport)

A distributed system is a collection of independent computers that are used jointly to perform a single task or to provide a single service.

Note 1: Characteristics

- Multiple computers
    CPU, memory, storage, I/O
- Interconnections
    variety of interconnection architectures
- Resources
    remote access to resources
    resource can be shared

1.1.2 Goals of Distributed Systems

- Transparency (hiding distribution)
    System presents itself as a single computer system
- Openness
    Interoperability, portability, heterogeneity
- Scalability
    Ability to grow

Transparency

Access: Hide differences in data representation and how a resource is accessed
Location: Hide where a resource is located
Migration: Hide that a resource may move to another location
Relocation: Hide that a resource may be moved to another location while in use
Replication: Hide that a resource is replicated
Concurrency: Hide that a resource may be shared by several competitive users
Failure: Hide the failure and recovery of a resource

Openness

- Interoperability
- Portability
- Heterogeneity
- Standard interfaces
- Interface Definition Language (IDL)

Scalability

' Three axis of scalability:
- Administratively
- Geographically
- Size (users, resources)

Algorithms vs Scalability Decentralized algorithms should be used:
- No machine has complete information about the system state
- Machines make decisions based only on local information
- Failure of one machine does not ruin the algorithm
- There is no implicit assumption that a global clock exists

Scaling Techniques
- Hiding communication latencies
    Asynchronous communication
    Client-side processing
- Distribution
    Split and spread functionality across the system
    Decentralize algorithms
- Replication (including caching)

If asynchronous communication cannot be used - communication should be reduced

1.1.3 Types of Distributed Systems

Distributed Computing Systems

- Cluster Computing Systems
  Just a bunch of computers all connected over a shared network
- Grid Computing Systems
  Layered System: Applications → Collective Layer → (Connectivity layer / Resource layer) → Fabric layer
- Cloud Computing
  Paradigm for enabling **network access** to a scalable and elastic pool of **shareable physical or virtual resources** with on-demand self-service provisioning and administration

Distributed Information Systems

- Transaction processing systems
  There are many information systems in which many distributed operations on (possibly distributed) data have to have the following behavior (either all of the operations are executed, or none of them is executed):
  BEGIN_TRANSACTION: Mark the start of a transaction
  END_TRANSACTION: Terminate the transaction and try to commit
  ABORT_TRANSACTION: Kill the transaction and restore the old values
  READ: Read data from a file, a table, or otherwise
  WRITE: Write data to a file, a table, or otherwise

Note 2: Distributed Transactions - Model

A transaction is a collection of operations that satisfies the following ACID properties:

Atomicity: All operations either succeed, or all of them fail. When the transaction fails, the state of the object will remain unaffected by the transaction.

Consistency: A transaction establishes a valid state transaction. This does not exclude the possibility of invalid, intermediate states during the transaction's execution.

Isolation (Serialisability): Concurrent transaction do not interfere with each other. It appears to each transaction T that other transactions occur either *before T*, or *after T*, but never both.

Durability: After the execution of a transaction, its effects are made permanent: changes to the state survive failures.

- Enterprise application integration
  Middleware as a communication facilitator in enterprise application integration
  Multiple applications communicate to the middleware which then talks to all the server-side applications

Distributed Pervasive Systems

Pervasive systems:
- Embedded devices
- Mobile devices
- Heterogeneous networks
- (Autonomic) Adaptation to context changes
  Adaptation to changes in the infrastructure
  Adaptation to user tasks/needs

Requirements for pervasive systems:
- Embrace contextual changes
- Encourage ad hoc composition
- Recognize sharing as the default

Home Systems (Smart Homes) Integration of entertainment and appliances into an "intelligent" adaptive system. May include health-monitoring and also provide support for independent living of the elderly.

Sensor Networks There is a variety of sensor networks, e.g.
- A small set of sensors supporting smart home

4

- A network of thousands of sensors providing climate monitoring

1.2 Architectures of Distributed Systems

Architecture styles

- Layered architectures
- Object-based architectures
- Data-centered architectures
- Event-based architectures

System architectures

(how software components are distributed on machines)

- Centralized architectures (client-server: two-tiered, three-tiered, N-tiered)
  - The simplest organization is to have only two types of machines:
    * A client machine containing only the programs implementing (part of) the user-interface level
    * A server machine containing the rest
      the programs implementing the processing and data level
- Decentralized architectures (peer-to-peer)
  - Overlay network is constructed in a random way
  - Each node has a list of members but the list is created in unstructured (random) way
- Hybrid architectures (edge-server, collaborative DS)
  Clients participate in providing services: e.g. file sharing, when part of file is downloaded it's seeded to other clients

Note 3: Application Layering

- The user-interface level
- The processing level
- The data level

1.2.1 Adaptability and self-management in DS

- Role of middleware is to provide some degree of distribution transparency
  Hiding distribution of data, processing and control

- Middleware may have a particular architectural style, e.g.
    Object-based (CORBA)
    Event-based (most of middleware built for adaptive, context-aware applications)
- Middleware should be adaptive to meet requirements of various applications

Using interceptors to adapt control

- Interceptors change flow of control and allow additional code to be executed
- Many object-based DS use interceptors to change flow of control in the object invocation
    Requests (object invocations) can be intercepted
    Messages can be intercepted

General Approaches to Adaptive Software

- **Separation of concerns**
    e.g. *aspect oriented programming* (not very successful)
- **Computational reflection**
    e.g. *reflective middleware*
- **Component-based design**
    Adaptation through composition
      Statically at design time
      Dynamically at run time (requires support for late binding)

Self-management / autonomic computing

- Many distributed systems are complex and require self-management and behaviour adaptation
- Autonomic computing (or *self-**)
  - Self-configuration
  - Self-healing
  - Self-management
  - Context-awareness, etc
- Adaptation takes place by one or more feedback control loops

The Feedback Control Model

Figure 1.1: The logical organisation of a feedback control system

1.2.2 Processes

> **Note 4: Core OS functionality**
>
> - Process Manager
>   Process lifecycle
> - Thread Manager
>   Thread lifecycle
> - Communication Manager
>   Communication between threads attached to different process
> - Memory Manager
>   Physical/virtual
> - Supervisor
>   Dispatching interrupts, traps, ...

Definitions

- **Process**
    - Program being executed
    - Consists of address space and execution state (program counter, stack pointer, processor status word, contents of registers and system-call state)
- **Thread** (of execution)
    - Abstraction of activity within process

Thread has less overhead and so interruptions and switching can be faster. Threads also have shared memory

The use of threads in DS

- Starting thread to handle incoming request much cheaper than starting new process
    - single-threaded server prevents simple scaling to multiprocessor
    - as with client: hide network latency by reacting to next request while previous being transmitted
- Better structure than nonblocking I/O

- single thread, blocking I/O doesn't scale up
- simple blocking calls simpler overall structure

> **Note 5: Multithreaded Servers**
>
> **Threads:** Parallelism, blocking system calls
> **Single-threaded process:** No parallelism, blocking system calls
> **Finite-state machine:** Parallelism, nonblocking system calls

Virtualisation in DS

Interface at different levels:
1. An interface between the hardware and software consisting of machine instructions
     that can be invoked by any program
2. An interface between the hardware and software, consisting of machine instructions
     that can be invoked only by privileged programs, such as an operating system
3. An interface consisting of system calls as offered by an operating system
4. An interface consisting of library calls
     generally forming what is known as an application programming interface (API)
     in many cases, system calls are hidden by an API

Clients/Servers

> **Note 6: Clients**
>
> - Clients are often **thin**
>   Large portion of processing required to support sophisticated user interfaces can by provided by servers
> - Clients should support transparencies
>   communication hiding
>   access transparency (client stub)
>   location, migration, relocation transparencies

Servers General Design Issues
Stateful server: maintains information on clients
     hard to handle failures
Stateless server: knows nothing about clients
     can change own state without informing clients

1.3 Communication

- Interprocess communication in DS is based on message exchange
- Low level (transport) message passing does not provide distribution transparency
  higher level models are needed
- Distributed application need a variety of communication semantics
  - Remote Procedure Call (RPC)
  - Message-Oriented Middleware (MOM)
  - Data streaming

1.3.1 Layered Protocols

Layers, interfaces, and protocols in the **Open Systems Interconnection Reference Model** (OSI)
Application: Application protocol
Presentation: Presentation protocol
Session: Session protocol
Transport: Transport protocol
Network: Network protocol
Data link: Data link protocol
Physical: Physical protocol

Middleware Layer

Provides common protocols that can be used by many different applications
- Security protocols
- Transaction protocols
- High-level communication models (e.g. RPC, message queuing services, etc)

1.3.2 Types of Communication

Synchronicity in Communication

> **Persistent communication** - message is stored in the system until receiver becomes active

Persistent asynchronous communication. A will send a message to B and not wait for response back

Persistent synchronous communication. A sends a message and wait's for B to respond before continuing execution

Transient asynchronous communication. A sends a message, B will only receive the message if it is running

Receipt-based transient synchronous communication. A sends a message, B receives the message if it is active and returns an acknowledgement

Delivery-based transient synchronous communication at message delivery. A sends a message and waits for the acknowledge from B, B will receive the message and send an acknowledge when it starts on the job

Response-based transient synchronous communication. A sends a message and waits for the acknowledge from B, B will receive the message and send an acknowledge when it has finished the job

Steps of a Remote Procedure Call

1. Client procedure calls client stub (procedure invocation)
2. Client stub builds message, calls local OS
3. Client's OS sends message to remote OS
4. Remote OS gives message to server stub
5. Server stub unpacks parameters, calls server
6. Server does work, returns result to the stub
7. Server stub packs it in message, calls local OS
8. Server's OS sends message to client's OS
9. Client's OS gives message to client stub
10. Stub unpacks result, returns to client

> **Note 7: Parameter Passing**
>
> **Parameter marshalling:** There's more than just wrapping parameters into a message. Client and Server:
> - have different data representations
> - have to agree on the same encoding (basic/complex data values)
> - Interpret data and transform them into machine-dependent representation

1.3.3 Parameter Specification and Stub Generation

- Parameters passed by value do not pose problems
- Parameters passed by reference - some partial solutions exist (e.g. a small array can be sent to the server)

- Interfaces (procedures) are often specified in Interface Definition Language (IDL) and compiled into stubs

Writing a Client and a Server

Three files output by the IDL compiler:
- A header file (e.g. interface.h, in C terms)
- The client stub
- The server stub

1.3.4 Binding a Client to a Server

- Registration of a server makes it possible for a client to locate the server and bind to it
- Server location is done in two steps:
    1. Locate the server's machine
    2. Locate the server on that machine

1.3.5 Message-Oriented communication

- RPC (and RMI) are **synchronous** and **transient** (both sender and receiver have to be active
- Transport level communication like TCP and UDP also **require that the sender and receiver are active** (transient communication)
- Some applications require that messages are kept in the system until the receiver becomes active (e-mail is one example, but there are many more applications of this type)

Asynchronous Communication

- MOM – Message-Oriented Middleware
- Can store messages in the system
- Applications communicate by inserting messages in queues
- Messages are delivered by an overlay network of application layer servers (routers)
- Each application has its own private queue to which other applications can send messages
- There is no guarantee on when the message will be delivered and that it will be read

Message-Queueing Model

Put: Append a message to a specific queue – non-blocking

Get: Block until the specified queue is nonempty, and remove the first message

Poll: Check a specified queue for messages, and remove the first. Never block

Notify: Install a handler to be called when a message is put into the specified queue

Message brokers

- Main role of brokers is to transform messages from sender's format to receiver's format
- This functionality can be generalised to brokers matching applications based on messages
    Applications *publish* messages
    Applications *subscribe* for messages

> **Note 8: MCA: Message Channel Agent**
>
> Responsible for checking a message, wrapping it, and sending it to associated receiving ends

Channels

Transport type: Determines the transport protocol to be used

FIFO delivery: Indicates that messages are to be delivered in the order they are sent

Message length: Maximum length of a single message

Setup retry count: Specifies maximum number of retries to start up the remote MCA

Delivery retries: Maximum times MCA will try to put received message into queue

Message Transfer

MQopen: Open a (possibly remote) queue

MQclose: Close a queue

MQput: Put a message into an opened queue

MQget: Get a message from a (local) queue

1.3.6 Transient Communication

Berkeley Sockets

Socket: Create a new communication end point

Bind: Attach a local address to a socket

Listen: Announce willingness to accept connections

Accept: Block caller until a connection request arrives

Connect: Actively attempt to establish a connection

Send: Send some data over the connection

Receive: Receive some data over the connection
Close: Release the connection

1.3.7 Actor model for communication

- The actor model adopts the philosophy that *everything is an actor*
- Can respond to a message it receives, can send a finite number of messages to other actors
- Enabling asynchronous communication and control structures as patterns of passing messages
- Resemble the Enterprise application integration framework

1.3.8 Stream-oriented communication

- Some applications need to exchange time-dependent information e.g. video, audio (continuous media)
- Previously discussed models do not consider time
- In continuous media the temporal relationship between different data items is essential

1.3.9 Different transmission modes

- Asynchronous transmission mode
  Sequential transmission without restrictions on when data is to be delivered (e.g. file transfer)
- Synchronous transmission mode
  Max end-to-end delay for each unit in a data stream (e.g. sensor sample temperature)
- Isochronous transmission mode (streams)
  Data transfer bounded by maximum and minimum end-to-end delay (bounded jitter e.g. audio/video)

1.3.10 Streams

Streams can be simple or complex (i.e. can include several related substreams)
- Unidirectional (source $\rightarrow$ sinks)
- Simple (single flow of data e.g. audio, video)
- Complex (stereo audio, movie)

Streams and QoS

Streams need timely delivery of data

- Non functional requirements (time, bandwidth, volume, reliability) are expressed as Quality of Service (QoS)
- There are many models for QoS specifications: e.g. IntServ, DiffServ, MPLS
- In IntServ (Integrated Services), QoS is specified as a flow specification – flow reservation
- In DiffServ (Differentiated Services), QoS is specified for a class (e.g. expedited forwarding class)

Properties for Quality of Service:
- The required bit rate at which data should be transported (application specific)
- The maximum delay until a session has been set up (when to start sending data)
- The maximum end-to-end delay
- The maximum delay variance, or jitter
- The maximum round-trip delay

Stream Synchronisation

- Given a complex stream, how are substreams synchronised?
- Synchronisation takes place at the level of data units (e.g. synchronise two streams only between data units)
- Issues which need to be considered:
  Mechanisms for synchronising streams
  Distribution of those mechanisms

Distribution of synchronisation mechanisms

- The receiving side needs to have a complete synchronisation specification
- Synchronisation specification can be multiplexed together with other substreams into a complex stream (and is demultiplexed after receiving)

Multicast communication

- Application level multicasting – setting a path for information dissemination
  - Nodes (applications) organise into an overlay network
  - Overlay network disseminates information
  - Network layer routing is independent of the overlay (communication may not be optimal)
- Overlay organisation
  - Nodes may organise themselves into a tree, or

– Nodes organise themselves into a mesh network

Gossip-based data dissemination

- Epidemic protocols are often used for data dissemination
  – There is no central component which co-ordinates dissemination
  – Information is propagated using local information only
- Node is *infected* if is has data which it wants to spread
- Node which has not seen this data is *susceptible*

> **Note 9: Information Dissemination Models**
>
> - Anti-entropy propagation model
>   – Node P picks another node Q at random
>   – Subsequently exchanges updates with Q
> - Approaches to exchanging updates
>   – P only pushes its own updates to Q
>   – P only pulls in new updates from Q
>   – P and Q send updates to each other
> - One variant is the "gossiping" protocol

1.4   Naming

1.4.1   Names, Identifiers and Addresses

- Names refer to entities
    e.g. processes, users, mailboxes, network connections
- Special types of names exist:
  – Addresses
  – Identifiers
  – Human-friendly names

Addresses

- To use (operate on) an entity you need an access point
- The name of an access point is an *address*
- Entities can have more than one access point at a time
- An entity may change access points e.g. when it changes location

- *Location-independent names* are not tied to an address

Identifiers

Special type of name that:
- Refers to at most one entity
- Always refers to the same entity
- Are limited to one per entity

Human-friendly names

Generally represented as a character string e.g. `www.news.com.au` or `/root/ryan/slides`

1.4.2   Flat Naming

Random bit strings (name doesn't help locate access point)
- Names must be resolved to address
- There are many name resolution approaches for flat names:
  – Broadcasting (or multicasting) approaches
  – Forwarding pointers
  – Home-based approaches (e.g. Mobile IP)
  – Distributed Hash Tables (DHTs)
  – Hierarchical approaches

Broadcasting

Broadcasting (or multicasting) the ID requesting the entity to return its current address. e.g. ARP (Address Resolution Protocol)

Forwarding Pointers

- When entity moves, it leaves a pointer to its new location
- To find entity, must follow trail of pointers to entity's current location
- Problems:
  – Can get long chains of pointers (not scalable)
  – If chain breaks, can't find entity
- Potential solution:
    Stub Scion Pair (SSP) chains

Home-Based Approach

- Approach designed for mobile entities

- Entities have home location that keeps track of entity's current location
- Example of home-based approach is Mobile IP:
  - Mobile entity has fixed IP address
  - All traffic to mobile entity goes to entity's Home Agent (home location)
  - Home Agent forwards traffic to Care-Of-Address (mobile entity's IP in its current network)
  - Whenever mobile entity changes network, it gets new Care-Of-Address, which it registers with Home Agent

Distributed Hash Tables

- Consist of many distributed nodes
- Map data to a key value using a hash function
- Each node is responsible for key values (and the associated data) in a particular range
- Have fast lookup times

Hierarchical Approach

- Network is divided into **domains**, with each domain subdivided into smaller subdomains
- **Leaf domain** is lowest level domain (typically LAN or Cell in mobile phone network)
- Each domain has a **directory node** that keeps track of entities in that domain
- Directory nodes for leaf domains store address of entities in that domain
- Directory nodes for non-leaf domains store reference to lower-level domain containing entities

1.4.3 Structured Naming

- Flat names are not convenient for humans
- Structured names:
  - Composed from simple human-readable names
  - Generally supported by naming systems
  - Names are organised into a **name space**

Name Spaces

- Name spaces for structured names are represented as a labelled, directed naming graph in which:
  - A **leaf node** represents a (named) entity

- A **directory node** is an entity that refers to other nodes
  - Outgoing edge is represented as (edge label, node identifier)
  - **Root node** has only outgoing edges
- Naming graphs are usually directed acyclic graphs
- Each path in a naming graph can be referred to by a sequence of edge labels separated by a special character

Name Resolution in a Name Space

- Need to know how and where to start → need closure mechanism
- Closure mechanism deals with selecting initial node in a name space from which name resolution is to start
- Closure mechanisms are often implicit

Linking and Mounting

- Aliases
  - Another name for the same entity
  - Symbolic links in UNIX enable more than one path for a file
- Mounting
  - Used to merge different name spaces transparently

Name Resolution

- Structured names resolved using naming service (implemented by name servers)
- Naming service allows addition, removal and lookup of names
- DS naming services are distributed
- Large-scale distributed naming services are usually hierarchical:
  - **Global level:** high-level directory nodes
  - **Administrational level:** mid-level directory nodes grouped into separate administrations
  - **Managerial level:** low-level directory nodes within a single administration
- The Domain Name System is a good example of a structured name resolution mechanism

DNS

- Large distributed name service used by Internet
- DNS name space is hierarchical

- DNS labels use alphanumeric character strings separated by a "."
- A path name in DNS is called a domain name
- DNS is primarily used to lookup IP addresses for hosts and mail servers
- Client (e.g. browser) contacts name resolver
- Name resolver uses either **iterative** or **recursive** approach
- Iterative name resolution:
  1. Name resolver contacts name servers for help
  2. Starts at root name server
  3. Each name server resolves as much of name as it can before referring name resolver to another name server who knows more
  4. Process repeats until name is fully resolved, or cannot be resolved anymore
- Recursive name resolution:
  1. Name resolver sends name to root name server
  2. Intermediate result not passed to client, rather it is sent to next name server
  3. Process repeats until name is fully resolved, then servers pass fully resolved name back
  4. Root name server passes fully resolved name to name resolver

Comparison between Iterative and Recursive:
- Recursive approach has higher performance demand on name servers
- Recursive approach can use result caching more effectively
- Iterative approach has higher communication costs

1.4.4   Attribute-based Naming

- Each entity is described by *(attribute, value)* pairs
- Attribute-based queries:
     Users specify attributes they are looking for
     Naming system should return one (or more) entities with the specified attributes
- Attribute-based naming systems are commonly known as **directory services**
- Lightweight Directory Access Protocol (LDAP) is a common directory service

LDAP

- Derived from OSI X.500 directory service

- LDAP directory service stores *directory entries*
- Directory entries consist of (attribute, value) pairs
- Directory Information Base (DIB) is collection of all directory entries in an LDAP service
- Each LDAP directory entry has globally unique name (Directory Information Tree) based on hierarchy of naming attributes
     e.g. /C=NL/O=Vrije/OU=Comp.Sc

Table 1.1: Simple Example of LDAP directory entry

Attribute	Abbr	Value
Country	C	NL
Locality	L	Amsterdam
Organisation	O	Vrije
OrganisationalUnit	OU	Comp.Sc
CommonName	CN	Main Server
Mail_Servers	–	137.37.20.3, 130.37.24.6
FTP_Server	–	130.37.20.20
WWW_Server	–	130.37.20.20

Decentralised Schemes

- Driven by advent of peer-to-peer
- Need efficient mapping of (attribute,value) pairs to avoid exhaustive search of networks
- Two Approachs:
  - Distributed Hash Tables (INS/Twine)
  - Semantic Overlay Networks

Note 10: Distributed Hash Tables

- In INS/TWINE each entity (resource) is described by a hierarchical attribute-value tree (AVTree)
- Every path from root of AVTree gets unique hash value
- Node in DHT responsible for hash value will keep reference to actual resource
- Query for (type-book) will get hashed to value 5 and sent to node responsible for storing hash value 5

- Every machine has a timer that generates an interrupt H times per second
- There is a clock in machine p that **ticks** on each timer interrupt. Denote the value of that clock by $C_p(t)$, where t is UTC time
- Ideally, we have that for each machine p, $C_p(t) = t$ or, in other words, $\frac{dC}{dt} = 1$

If $\frac{dC}{dt} > 1$, fast clock. If $\frac{dC}{dt} = 1$, perfect clock. If $\frac{dC}{dt} < 1$, slow clock.

1.5 Synchronisation

1.5.1 Clock Synchronisation

- The problem with clocks in DS
- Physical clocks
- clock synchronisation algorithms:
    network time protocol (NTP)
    the berkeley algorithm
    clock synchronisation in wireless networks

Physical Clocks

- Sometimes we need the exact time in DS
- **Universal Coordinated Time (UTC)**
    - Based on the number of transactions per second of the Cesium 133 atom
    - At present, the real time is taken as the average of approx. 50 Cesium clocks around the world (International Atomic Time - TAI)
    - Introduces a leap second from time to time to compensate that days are getting longer
- UTC is broadcast through short wave radio and by satellite. Satellites can give an accuracy of about $\pm$ 0.5ms

Clock Synchronisation Algorithms

- Suppose we have a distributed system with a UTC receiver somewhere in it we still have to distribute its time to each machine
- Assumptions:

1.5.2 Logical Clocks

- The order in which events occur in the DS is often more important than the time that they occurred
- The order of events can be established using logical clocks

Lamport's Algorithm

- Events in the DS can be ordered using Lamport's **happened-before** relation
- The **happened-before** relation on the set of events in a DS is the smallest relation satisfying:
  - if a and b are events in the same process, and a occurs before b, then $a \rightarrow b$ is true
  - if a is the event of a message being sent by one process, and b is the receipt of that message in another process, then $a \rightarrow b$ is also true
- This relation is transitive:
  - $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$
- In Lamport's logical clock algorithm:
  - each event a has an associated time $C(a)$ based on the local clock
  - between any two events, the clock must tick at least once (i.e. no two events ever occur at the same time)
  - messages carry their sending time according to the sender's clock e.g. $C(b)$

  - when a message is received, its time is compared against the local clock. If the local clock is less than $C(b)$, **it is set to** $C(b) + 1$
- Events occurring in processes that do not interact (even indirectly through third parties) are said to be concurrent
- Nothing can be said about the order of concurrent events
- For example:
  - Events x and y occur in two different processes (that do not interact at all)
  - The happened-before relation cannot be applied as x and y are concurrent
  - This means that $x \rightarrow y$ is not true, and $y \rightarrow x$ is not true

Vector Clocks

- Lamport's algorithm ensures that if the happened-before relation exists between events a and b, then $C(a) < C(b)$
- But if $C(c) < C(d)$, Lamport's algorithm doesn't guarantee that c happened before d
- The problem is that Lamport's algorithm does not capture **causality**
- Causality can be captured using vector clocks
- The vector clock for an event a is signified by $VC(a)$
- For two events, a and b, if $VC(a) < VC(b)$, then event a causally precedes event b
- To construct a vector clock, each process P_i maintains a vector VC_i
  - $VC_i[i]$ is the value of the logical clock at P_i
  - If $VC_i[j] = k$, P_i knows that at least k events have occurred at P_j
- For every event that occurs at P_i the vector value $PC_i[i]$ is incremented by one
- A process' vector is piggybacked onto all messages sent by that process
- Every time a message is received, the recipient process updates its own vector (VC_r)
- If we assume the message vector is VC_m, the update is performed by setting $VC_r[k] = MAX(VC_r[k], VC_m[k])$ for each k
- The issue of total ordering and causal ordering of messages by the communication system is controversial
- Total or causal ordering can also be provided in the application (**end-to-end** argument)

1.5.3 Mutual Exclusion

- Processes in a distributed system may want exclusive access to a shared resource
- A mutual exclusion mechanism is required to prevent corruption (or inconsistent updates) of that resource
- How to achieve mutual exclusion in DS?

A Centralised Algorithm

- Coordinator process enforces mutual exclusion over resource
- Processes must ask coordinator for permission to access resource
- Benefits of centralised approach:
  - easy to implement
  - low message overhead
  - fair (access requests are processed in order)
- Drawbacks:
  - coordinator is a single point of failure
  - coordinator can be performance bottleneck

Distributed, with no topology imposed (by Ricart and Agrawala)

- A process wanting to access a shared resource sends a message to all other processes. The message contains:
  - the requested resource's name
  - the requesting process' process id
  - the requesting process' logical time
- Recipients of the message follow one of three behaviours. If the recipient:
  - doesn't want the resource, it sends back OK
  - currently holds the resource, then it checks if the logical time in the message is less than its own logical time. If so, it sends back OK. If not, it queues the message and sends back nothing.
- To access the shared resource, a process must receive an OK from all other processes
- When a process is finished with a resource it:
  - sends OK messages to processes in its queue
  - deletes its queue
- Benefits of approach:
  - solution is fair
  - does not need a single coordinator
- Drawbacks
  - all processes are involved in all decisions (one slow process slows down others)
  - large number of messages required
  - single point of failure replaced with n points of failure
  - processes must have accurate group membership list

Distributed, using a ring topology

- Uses a logical ring to order processes
- Processes can only access the shared resource while in possession of a token
- The token is passed on to the next node in the ring if:
  - the current token holder does not want to access the shared resource
  - if the token holder is finished accessing the shared resource
- The token circulates around the ring in one direction
- Benefits of this approach:
  - algorithm is simple
- Drawbacks:
  - token must be regenerated if lost
  - crashed processes can stop circulation of token
  - potentially have to wait for every other process to use token before it is your turn

Table 1.2: Comparison of Algorithms

Algorithm	Messages per entry/exit	Delay before entry (in message times)	Problems
Centralised	3	2	Coordinator crash
Distributed	$2(n-1)$	$2(n-1)$	Crash of any process
Token ring	1 to ∞	0 to $n-1$	Lost token, process crash

1.5.4 Election Algorithms

- Many distributed algorithms require that one of the processes acts as a coordinator
- An election is used to dynamically select co-ordinator
- Election algorithms needed so that at the end of the election all processes agree on coordinator
- Common election algorithms are:
  - the Bully Algorthm
  - a Ring Algorithm
- Different algorithms needed for:
  - Wireless network environments
  - Large-scale distributed systems

The Bully Algorithm

- When process P notices coordinator is non-responsive, it initiates an election
- Election conducted as follows:
  1. P sends an ELECTION message to all processes with higher process number
  2. If no one responds, P wins the election and becomes coordinator
  3. If one of the higher-ups answers, it takes over. P's job is done

A Ring Algorithm

- Assumes processes are logically ordered (each node knows its successor in ring)
- When process P notices coordinator is non-responsive, it initiates election
- Election conducted as follows:
  - P sends ELECTION message (containing P's process num) to successor
  - Recipients add own process num to message and pass to their successor
  - Message gets back to P, who changes message type to COORDINATOR
  - COORDINATOR message circles ring again
  - Process with highest process num in COORDINATOR message becomes co-ordinator

for Wireless Environments

- Previously described election algorithms need:
  - reliable message passing
  - stable network topology

- These aren't always present in wireless environments
- The following election protocol for wireless environments attempts to overcome these problems
- Node that calls election becomes **source node**
- Source node sends ELECTION message to all neighbours
- The first time a node receives ELECTION message it:
  - marks sender as parent
  - forwards ELECTION message to all its neighbours
- Subsequent ELECTION messages (not from parent) are acknowledged only
- Nodes wait a set time for acknowledgements from neighbours, before sending own acknowledgement to parent
- Acknowledgements contain information on the resource capacities of the node's best neighbour (e.g. battery power)
- The source node uses the acknowledgement information to select the coordinator

for Large-Scale Systems

- Previous algorithms select one node only
- Large-scale systems may require many local coordinators (e.g. peer-to-peer networks superpeers keep index of content on neighbours to speed searches)
- Superpeers should:
  - offer regular nodes low-latency access
  - be evenly distributed throughout network
  - exist in predefined proportion to regular nodes
  - serve no more than a set number of regular nodes
- Two approaches for selecting superpeers:
  - using a Distributed Hash Table (DHT) identifier
  - using repulsion forces
- Distributed Hash Table identifier:
  - fraction of DHT identifier space is reserved for superpeers
  - reserve the first (i.e. leftmost) k bits to identify superpeers
  - need N superpeers: use *ceil(log2(N))* bits of any **key** to identify these nodes
- Repulsion forces approach:
  - n tokens spread across peer-to-peer overlay

- **each node holding a token** learns about other token-holders
- each token is "repulsed" by nearby tokens (token holder sends token to another peer if too many tokens nearby)
- tokens passed around network until tokens spread evenly across network
- token must be held by a node for a set time period before node can become superpeer

1.6 Consistency and Replication

- Replication of services
  - Increase availability
  - Enhance reliability (switch to another on crash)
  - Improve performance (load balance, data closer, support scaling in numbers and area)
- Management of replicated data:
  - Scalability of keeping replicas consistent

1.6.1 Data-centric consistency

Data store: distributed collection of storages accessible to clients

Consistency model: contract between (distributed) data store and processes, in which data store specifies precisely what results of read and write operations are in presence of concurrency

Strong consistency: operations on shared data *coherent*
- Strict consistency (related to time)
- Sequential consistency (what we are used to)
- Causal consistency (maintains only causal relations)

Weak consistency: coherence only when shared data locked, unlocked
- Entry consistency

Strict Consistency: Any read to a shared data item X returns the value stored by the **most recent** write operation on X

Sequential Consistency: Result of any execution same as if operations of all processes executed in some sequential order, and operations of each individual process appear in this sequence in order specified by its program

Causal Consistency: Writes that are potentially **causally** related must be seen by all processes in same order. Concurrent writes may be seen in different order on different machines

> the weaker the consistency model, the easier to scale

Grouping Operations - Weaker Consistency

- Most applications use transactions or enter critical sections
- Consistency of whole set of operations is of interest not single read/writes
  - During transactions or operations in critical sections (CS) concurrent access to data used in CS is not allowed
  - Entering and leaving CS can be modelled by shared synchronisation variables
    * When process enters CS it should acquire synchronisation variables
    * When it leaves it releases these variables

Note 15: Summary of Consistency Models

Consistency models not using synchronisation variables
Strict: Absolute time ordering of all shared accesses
Sequential: All processes see all shared accesses in the same order. Accesses are not ordered in time
Causal: All processes see causally-related shared accesses in the same order
Models with synchronisation variables
Entry: Shared data pertaining to a critical region are made consistent when a critical region is entered

Eventual Consistency

For very large databases
- with tolerance for high degree of inconsistency
- updates guaranteed to propagate to all replicas
- if no updates for long time, all replicas gradually become consistent (identical copies)

Eventual consistency works well if always the same replica is accessed

1.6.2 Client-centric consistency

Goal: avoid system-wide consistency. Concentrate on single client's needs.
- Mobile clients
- Clients may have a distributed store but at a time work only on a local replica
  - No simultaneous updates
  - Most operations involve reading data

Monotonic Reads: If process reads value of data item x, any successive read operation on x by that process will always return that same or more recent value

1.6.3 Replica Management

Replica Placement

Model: consider objects (don't care if data or code)
Distinguish different processes: process capable of hosting replica of object or data

- **Permanent replicas:** process/machine always has replica
- **Server-initiated replica:** Process can dynamically host replica on request of server in data store (push cache)
  - Used to improve performance
- **Client-initiated replica:** Process can dynamically host replica on request of client
  - Client caches

Update Propagation

There are three possibilities:
- Propagate only a notification of an update
  - invalidation protocols
- Transfer data from one copy to another
  - used for high read-to-write ratio
- Propagate the update operation to other copies
  - active replication

Figure 1.2: Pull versus Push Protocols for updates

Issue	Push-based (server-based protocols)	Pull-based (client-based protocols)
State at server	List of client replicas and caches	None
Messages sent	Update (and possibly fetch update later)	Poll and update
Response time at client	Immediate (or fetch-update time)	Fetch-update time

1.6.4 Consistency protocols

Primary-Based Protocols: Fixed server to which all read and write operations forwarded
Replicated-write protocols:
Quorum-based protocols: Ensure that each operation carried out in such a way that majority vote established: distinguish **read quorum** and **write quorum** (Gifford's Scheme)

> **Note 16: Implications of Gifford's Scheme**
>
> General rule (N_R read quorum, N_W write quorum):
> - get N_R replicas to agree before read, N_W for write
> - meet conditions:
> $N_R + N_W > N$ (prevents read-write conflicts)
> $N_W > \frac{N}{2}$ (prevents write-write conflicts)

Chapter 2

Tutorials

2.1 Introduction to Distributed Systems

What is the role of middleware in a distributed system? Middleware provides distributed transparency such as:
- Access Transparency
- Location Transparency
- Concurrency Transparency

Explain what is meant by distribution transparency and give examples of different types of transparency. Distribution transparency allows aspects of distributed systems (such as accessing of data or individual software components) to be hidden and appear to the end user as a single system. Examples:
- Access Transparency
- Location Transparency
- Migration Transparency
- Relocation Transparency
- Replication Transparency
- Concurrency Transparency
- Failure Transparency

Why is it sometimes so hard to hide the occurrence and recovery from failures in a distributed system? It can be difficult to identify the state of remote components. For example how do you tell the difference between unavailable resource and a slow resource?

Why is it not always a good idea to aim at implementing the highest degree of transparency possible? Aiming at the highest degree of transparency may lead to a considerable loss of performance that users are not willing to accept.

What is an open distributed system and what benefits does openness provide? An open distributed system offers services according to a clearly defined set of rules and interfaces. An open system is capable of easily interoperating with other open systems but also allows applications to be easily ported between different implementations of the same system. Furthermore, an open distributed system allows software/hardware components of different natures (e.g. Windows, Linux workstations etc) to participate in the system. A few benefits are:
- The same system can deliver the service to different types of clients and applications
- The same system can work given different computing environments
- The same system can be extended to allow computing and storage technologies by different vendors

Describe precisely what is meant by a scalable system. A system is scalable with respect to either its number of components, geographical size, or number and size of administrative domains, if it can grow in one or more of these dimensions without an unacceptable loss of performance.

Scalability can be achieved applying different techniques. What are these techniques? Scaling can be achieved through:
- Workload and data distribution. This requires distributed/parallel algorithms
- Using decentralised architecture
- Data replication, and caching

2.1.1 Architecture of Distributed Systems

If a client and a server are placed far apart, we may see network latency dominating overall performance. How can we tackle this problem?
1. Buffering the communication so there is enough data presented to the client while more data is being transferred
2. Perform more client-size processing and caching to reduce communication load
3. Multithreaded communication requests on client and server to reduce network latency

What is a three-tiered client-server architecture? A three-tiered client-server architecture consists of three logical layers, where each layer is, in principle, implemented at a separate machine.

The highest layer consists of a client user interface, the middle layer contains the actual application, and the lowest layer implements the data that are being used.

What is the different between a vertical distribution and a horizontal distribution?

Vertical distribution: Multiple layers, each implemented on a different machine

Horizontal distribution: Single layer, implemented across multiple machines (distributed database)

In a structured overlay network, messages are routed according to the topology of the overlay. What is an important disadvantage of this approach? When a message is routed across a structure overlay network (which is a logical network) the shortest path between source and destination may not be the physical shortest path. While the source and receivers may be logically very close to each other, they could be physically at the remotest part of the network.